A PROPOSED DEPLOYMENT · ON YOUR FLOOR, IN YOUR CLOUD, OR BOTH · IN PHASES

CYAN

How Cyan would be deployed: the work travels to your footage, the record travels everywhere.

Cyan runs beside your storage, on your premises or in the cloud. Your footage stays put; the record is yours, on your rack or in your cloud account; only small things cross between rooms, and masters are delivered from the site that made them.

SCAN · KEEP THIS PAGE

Two ways to run Cyan, and what travels between the rooms

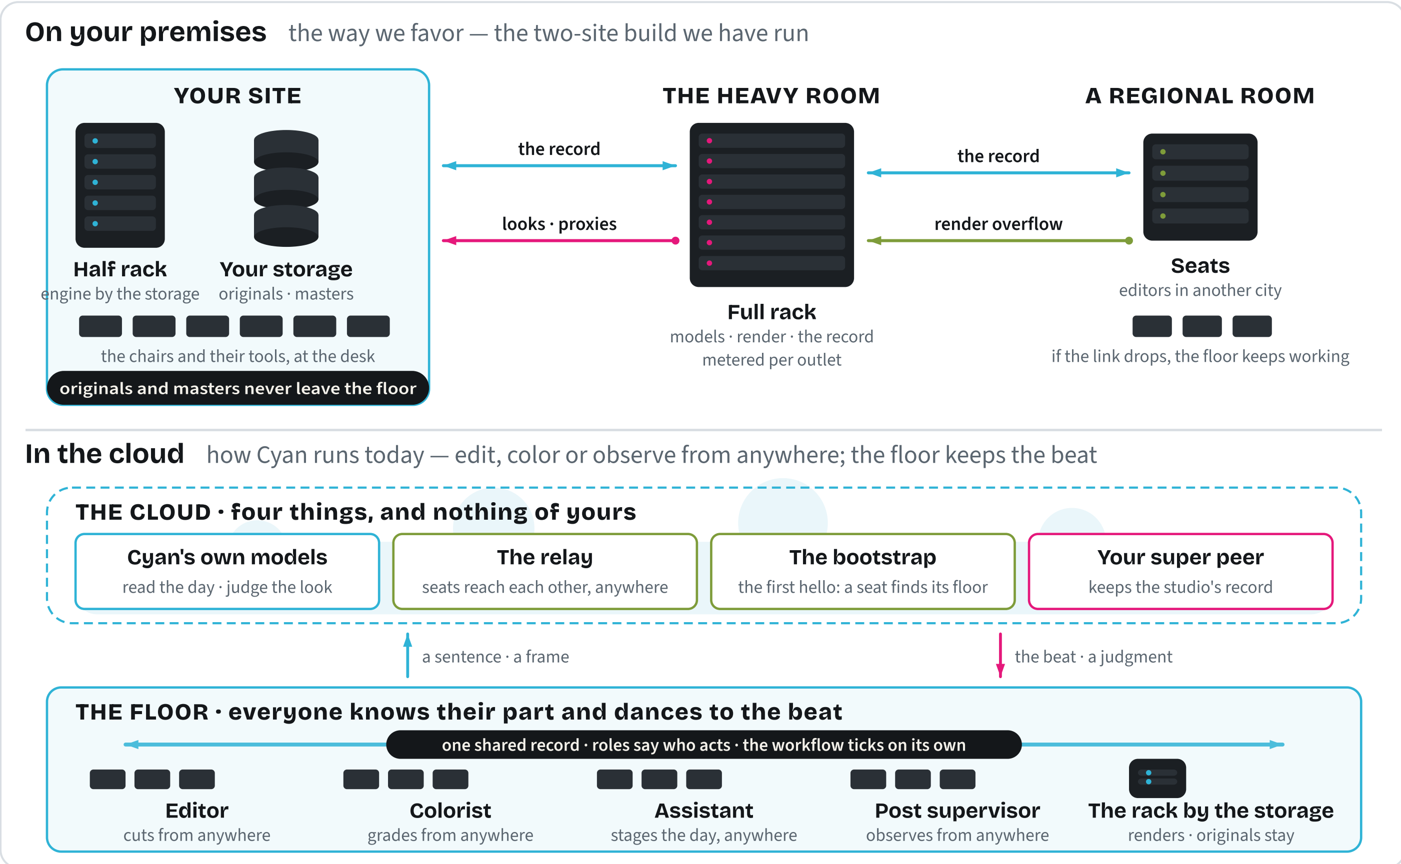

A phased build-out, as proposed

PHASE	WHERE	WHAT WOULD RUN	SCALE, AS AN ESTIMATE
One	Your site: half a rack, or one cloud room	The engine and its database on one machine, one reading model, working storage	One or two engines, single-digit seats, about five kilowatts
Two	The heavy room: a full rack, or a cloud region	GPU render workers, larger judging models, the activity log and usage metering, the hot cache	About sixteen GPUs within a ten-kilowatt rack, repeatable per rack
Three	A regional site, or a cloud room	Regional seats with their own permissions, overflow render, an optional shared model gateway	Seats and burst render; the main floor is never dependent on it

A second studio is the same playbook beside its own media. A rack grows when sustained GPU use passes seventy percent; a cloud room grows by the hour. Either way, every step carries its time and its cost.

What we need from you to start

- Where the media root lives, and how it is shared.
- A rack slot with power, or a cloud account, for the first room.
- The seats' machines and their DCC licenses: Premiere, Resolve, Pro Tools.

What you get in the first week

- A show set up from one sentence, on your floor.
- One shoot day run end to end: arrival, dailies, rough cut, lock, hand-offs.
- Every step receipted and metered on your own record.